

Lecture 7: Neural Nets and the Learning Function

Agenda

- Construction of Neural Nets
- Distance Matrices

Source: Sections VII.1 and IV.10 in Linear Algebra and Learning from Data (2019) by Gilbert Strang

Construction of Neural Nets

Learning Function $F(X, V)$ where x are the weights, and v are the feature vectors, the sample feature vectors (training dataset)

So those feature vectors V_0 come from the training data, either one at a time, if we're using stochastic gradient descent (discussed detail later) with mini-batch size 1

Or B at a time, if we're doing mini-batch of size B , or the whole thing, a whole epoch at once, if we're doing full-scale gradient vector

So those are the feature vectors, and these are the numbers in the linear steps, the weights

Weight matrix A_k multiply by V and bias vectors b_k that adds on to shift the origin. Optimize X and V

Take first step of learning Function F :

$v_1 = \text{ReLU}(F(A_1, b_1, v_0)) \rightarrow \text{non-Linear Step}$

ReLU (Rectified Linear Unit) is an activation function.

Construction of Neural Nets

Deep learning is **Continuous Piecewise Linear (CPL) functions**.

Linear for simplicity, **continuous** to model an unknown but reasonable rule, and **piecewise** to achieve the nonlinearity that is an absolute requirement for real images and data

Here is a first construction of a piecewise linear function of the data vector v .

Choose a matrix A_1 and vector b_1 .

Then set to zero (this is the nonlinear step) all negative components of $A_1 v + b_1$.

Then multiply by a matrix A_2 to produce 10 outputs in $w = F(v) = A_2(A_1 v + b_1) + \dots + A_{10}(A_{10} v + b_{10})$

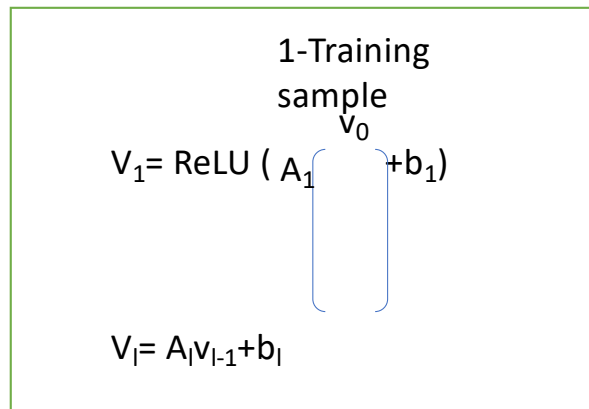
That vector $(A_1 v + b_1) + \dots$ forms a "hidden layer" between the input v and the output w .

Construction of Neural Nets

$(A_1 v_0 + b_1) \rightarrow$ linear step

$V_1 = \text{ReLU}(A_1 v_0 + b_1)$ non-linear step

Generally, $V_k = \text{ReLU}(A_k v_{k-1} + b_k)$ where $k=1, \dots, l$ (l : layers)



Don't do ReLU at the last layer, so it's just $A_l v_{l-1} + b_l$.

May not do a bias vector also at that layer, but you might, and this is the finally the output.

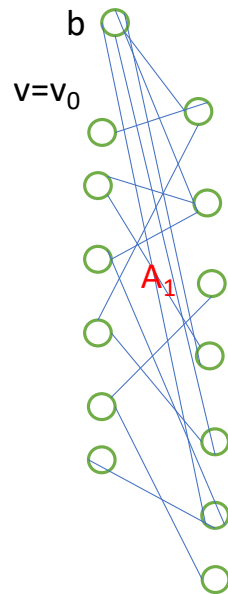
So this picture is clearer to distinguish between the weights

$x = A_1, b_1, A_2, b_2, \dots, A_l, b_l$ x really stands for all the weights that we compute up to A_l, b_l , so that's a collection of all the weights

Often weights x 's are undetermined because the number of X 's in A 's, b 's are greater than the number of v 's in the training sets

Construction of Neural Nets

1st layer: $v_1 = A_1 v_0 + b_1$



Choose x to min Loss Function $L = \left(\frac{1}{N}\right) \left[\sum_{i=1}^N F(x, v_i)\right]$

$$L(x) = \left(\frac{1}{N}\right) \left[\sum_{i=1}^N F(x, v_i) - \text{true}_i\right]$$

Question: do we use the whole function L at each iteration, or do we just pick only b , of the samples to look at iteration number K ?

So this is the $L(x)$ then added up over all v 's. This is what the neural net produces. It's supposed to be close to the true

Popular Loss functions:

(1) Square loss = sum of $\| \cdot \|_2^2 \rightarrow$ regression

(2) L^1 loss = sum of $\| \cdot \|_1 \rightarrow$ Lasso

(3) Hinge loss (-1,1 Classification)

(4) Cross-entropy loss (neural nets)

Distance Matrices

Question: distances squared:: $\|x_i - x_j\|^2 = d_{ij}$. Find positions x_j in \mathbb{R}^d (also find d i.e, the dimension of the space).

$\|x_i - x_j\|^2 = \text{given } d_{ij}$. Find x 's. Given $D = \{d_{ij}\}$ distances matrix, to find X matrix which gives the positions

In machine learning, you're given a whole lot of points in space, feature vectors (points) in a high-dimensional space, and those are related i.e., they are connected.

They tend to fit on a surface in high-dimensional space, a low- dimensional surface in high-dimensional space

Let's recognize the connection between distances and positions:

$$D = d_{ij} = \|x_i - x_j\|^2 = (x_i - x_j)^T (x_i - x_j) = x_i^T x_i - x_i^T x_j - x_j^T x_i + x_j^T x_j \quad (\text{entries in } D) \quad (1)$$

$x_i^T x_i$ produces a matrix with constant rows (no dependence on j). $x_j^T x_j$ produces a matrix with constant columns (no dependence on i) (For detail See Appendix foil 12)

$\|x_i\|^2$ and $\|x_j\|^2$ in both of those matrices are on the main diagonal of $G = X^T X$

Those are the numbers in the column vector $\text{diag}(G)$ (For detail See Appendix foil 12)

Distance Matrices

Middle terms $-2x_i^T x_j$ in (1) in last foil, $-2G = -2X^T X$

Rewrite (1) as an equation for the matrix D , using the symbol $\mathbf{1}$ for the **column vector of n ones**

That gives constant columns and $\mathbf{1}^T$ gives constant rows

So,
$$D = \mathbf{1} \text{diag}(G)^T - 2G + \text{diag}(G) \mathbf{1}^T \quad (2)$$

(Note: deduction of equation 2 is given in the next foil and For detail See Appendix in last foil)

Given D Find X // actually find $X^T X = G$ then find X from G

We'll find $X^T X$

Because we have dot products of X 's. Find out **what $x_i \cdot x_j$ is.**

Let's call this **matrix G for the dot product matrix**, and then find X from G .

Now let us say diagonal matrix $D_{ii} = (x_i, x_i)$. Let us write an equation for **G with dot matrix $X^T X$**

Distance Matrices

Solve equation (2) (from last foil) for $G = X^T X$

Place first point at the origin : $x_1 = 0$. For every $\|x_i - x_1\|^2$ is $\|x_i\|^2$

First column d_i of D (which is given) is the same as

$$\text{diag}(X^T X) = \text{diag}(G) = (\|x_1\|^2, \|x_2\|^2, \dots, \|x_n\|^2) \rightarrow \text{diag}(G) = d_1 \text{ and } \text{diag}(G) \mathbf{1}^T = d_1 \mathbf{1}^T$$

$$X^T X = G = -\frac{1}{2} D + \frac{1}{2} \begin{pmatrix} 1 \\ 1 \\ 1 \end{pmatrix} \begin{pmatrix} d_1 & d_2 & d_3 \end{pmatrix}^T + \frac{1}{2} \begin{pmatrix} d_1 \\ d_2 \\ d_3 \end{pmatrix} \begin{pmatrix} 1 & 1 & 1 \end{pmatrix}^T$$

Every colm Every row

Now G comes from D . G is positive semidefinite provided the distances in D satisfy the triangle inequality

Ref: Menger: *Amer. J. Math.* 53; Schoenberg: *Annals Math.* 36

This is the key equation

Matrix form

$$D = \begin{pmatrix} 1 \\ 1 \\ 1 \end{pmatrix} \begin{pmatrix} d_1 \\ d_2 \\ d_3 \end{pmatrix}^T + \begin{pmatrix} d_1 \\ d_2 \\ d_3 \end{pmatrix} \begin{pmatrix} 1 & 1 & 1 \end{pmatrix}^T - 2XX^T$$

$$XX^T = \frac{1}{2} \left[D - \begin{pmatrix} 1 \\ 1 \\ 1 \end{pmatrix} \begin{pmatrix} d_1 \\ : \\ d_i \end{pmatrix}^T - \begin{pmatrix} d_1 \\ : \\ d_i \end{pmatrix} \begin{pmatrix} 1 & 1 & 1 \end{pmatrix}^T \right]$$

Distance Matrices

Given XX^T Find X ($n \times n$)

Two leading candidates

- (1) Evaluate of $XX^T = Q\Lambda Q^T$
- (2) Elimination on XX^T

Find X up to an orthogonal transformation, as XX^T is symmetric

XX^T is positive or semidefinite, this is semidefinite. Given a semidefinite matrix and find a square root. Matrix is the XX^T and find X

There are many candidates, because if you find one i.e., any QX .
Because $Q^T Q$ is the identity.

Distance Matrices

(1) Take $X = Q\sqrt{\Lambda}Q^T = X^T$

$XX^T = (Q\sqrt{\Lambda}Q^T)(Q^T\sqrt{\Lambda}Q) = \Lambda = I = \text{identity matrix}$

(2) Elimination on $XX^T = LDU$ (L, a lower triangular, times D, the pivots, times U, the upper triangle)
= LDL^T (U is replaced by L^T)

Then $X = \sqrt{D}L^T$ (This is the Cholesky Factorization)

(Note: when X^TX , then X^TX coming correctly. X^T will be L^T . Transpose will give L. Square root of D will be \sqrt{D} . We'll give the D, and then the L^T is right)

Appendix

$$D = d_{12} = ||x_1 - x_2||^2 = (x_1 - x_2)^2 = x_1^2 - 2x_1x_2 - x_2^2 = x_1^T \cdot x_1 - 2x_1^T x_2 + x_2^T \cdot x_2$$

$$x_1^T \cdot x_1 = \begin{pmatrix} x_1 \\ 0 \end{pmatrix} (x_1 \quad 0) \text{ and } 2x_1^T \cdot x_2 = 2 \begin{pmatrix} x_1 \\ 0 \end{pmatrix} (x_2 \quad 0) \text{ and } x_2^T \cdot x_2 = \begin{pmatrix} x_2 \\ 0 \end{pmatrix} (x_2 \quad 0)$$

$$\begin{pmatrix} x_1 \\ 0 \end{pmatrix} (x_1 \quad 0) = \mathbf{1} \mathbf{diag} \left[\begin{pmatrix} x_1 \\ 0 \end{pmatrix} (x_1 \quad 0) \right]^T = \mathbf{1} \begin{pmatrix} x_1 \cdot x_1 \\ 0 \end{pmatrix}^T$$

Similarly,

$$D = d_{12} = \begin{pmatrix} 1 \\ 0 \end{pmatrix} \mathbf{diag} \begin{pmatrix} x_1^2 \\ 0 \end{pmatrix}^T - 2 \begin{pmatrix} x_1 \\ 0 \end{pmatrix} (x_2 \quad 0) + \mathbf{diag} \begin{pmatrix} x_2^2 \\ 0 \end{pmatrix} \begin{pmatrix} 1 \\ 0 \end{pmatrix}^T$$

$$D = \mathbf{1} \mathbf{diag}(G)^T - 2G + \mathbf{diag}(G) \mathbf{1}^T \text{ where } G = x_1^T \cdot x_2$$